CDP Production Deployment

****
CDP production deployment
****

-----*** Create a VPC ***-----

Create a elastic ip

Create VPC with a public and a private subnet

Create two security groups

-> In default security group

All traffic your-ip/32
All traffic private-subnet-cidr/24

-> in second security group (allowing communication only from edge)

All traffic default-sec-group-id
All traffic private-subnet-cidr/24

-----*** Launch public nodes ***-----

-----*** Update the server ***-----
$ sudo apt-get update && sudo apt-get dist-upgrade -y

-----*** Install NTP ***-----

$ sudo apt-get install ntp -y
$ sudo service ntp status
If it isn't running
$ sudo service ntp start

-----*** Set Swappiness ***-----

```
sudo sysctl -a | grep vm.swappiness
sudo sysctl vm.swappiness=1
echo 'vm.swappiness=1' | sudo tee --append /etc/sysctl.conf
```

```
$ sudo nano /etc/init.d/disable-transparent-hugepages
```

```
#!/bin/sh
### BEGIN INIT INFO
# Provides:          disable-transparent-hugepages
# Required-Start:    $local_fs
# Required-Stop:
# Default-Start:     2 3 4 5
# Default-Stop:      0 1 6
# Short-Description: Disable Linux transparent huge pages
# Description:       Disable Linux transparent huge pages, to improve
#                    database performance.
### END INIT INFO

case $1 in
  start)
    if [ -d /sys/kernel/mm/transparent_hugepage ]; then
      thp_path=/sys/kernel/mm/transparent_hugepage
    elif [ -d /sys/kernel/mm/redhat_transparent_hugepage ]; then
      thp_path=/sys/kernel/mm/redhat_transparent_hugepage
    else
      return 0
    fi

    echo 'never' > ${thp_path}/enabled
    echo 'never' > ${thp_path}/defrag

    unset thp_path
    ;;
esac
```

```
$ sudo chmod 755 /etc/init.d/disable-transparent-hugepages
```

```
$ sudo update-rc.d disable-transparent-hugepages defaults
```

>>Restart server

-----*** Setting root reserved space ***-----

$ sudo lsblk
$ sudo tune2fs -l /dev/nvme0n1p1
$ sudo tune2fs -l /dev/nvme0n1p1 | egrep "Block count|Reserved block count"
$ sudo tune2fs -m 1 /dev/nvme0n1p1
$ sudo tune2fs -l /dev/nvme0n1p1 | egrep "Block count|Reserved block count"

Note : Now, save the instance to an image, call it “Cloudera Manager” Make sure to check “No reboot”

-----*** Launch cluster nodes ***----

Launch reamining instances in private subnet

Select second security group for these instances

Preparing external database for cloudera

-----*** Install MySQL ***-----

sudo apt-get install mysql-server -y

-----*** Stop mysql ***-----

sudo service mysql stop

sudo service mysql status

-----*** Move innodb log file ***-----

mkdir mysqlbup

sudo su

cd /var/lib/mysql

mv ib_logfile0 /home/ubuntu/mysqlbup/

mv ib_logfile1 /home/ubuntu/mysqlbup/

exit

-----*** Edit the option file mysqld.cnf to cloudera recommended settings ***-----

cd /etc/mysql/mysql.conf.d/

sudo nano mysqld.cnf

==========

[mysqld]
datadir=/var/lib/mysql
socket=/var/lib/mysql/mysql.sock
transaction-isolation = READ-COMMITTED
Disabling symbolic-links is recommended to prevent assorted security risks;
to do so, uncomment this line:
symbolic-links = 0

key_buffer_size = 32M
max_allowed_packet = 16M
thread_stack = 256K
thread_cache_size = 64
query_cache_limit = 8M
query_cache_size = 64M
query_cache_type = 1

max_connections = 550
#expire_logs_days = 10
#max_binlog_size = 100M

#log_bin should be on a disk with enough free space.
#Replace '/var/lib/mysql/mysql_binary_log' with an appropriate path for your
#system and chown the specified folder to the mysql user.
log_bin=/var/lib/mysql/mysql_binary_log

#In later versions of MySQL, if you enable the binary log and do not set
#a server_id, MySQL will not start. The server_id must be unique within
#the replicating group.
server_id=1

binlog_format = mixed

read_buffer_size = 2M
read_rnd_buffer_size = 16M
sort_buffer_size = 8M

```
join_buffer_size = 8M

# InnoDB settings
innodb_file_per_table = 1
innodb_flush_log_at_trx_commit  = 2
innodb_log_buffer_size = 64M
innodb_buffer_pool_size = 4G
innodb_thread_concurrency = 8
innodb_flush_method = O_DIRECT
innodb_log_file_size = 512M

[mysqld_safe]
log-error=/var/log/mysqld.log
pid-file=/var/run/mysqld/mysqld.pid

sql_mode=STRICT_ALL_TABLES
```

==========

-----*** Make mysql start on boot ***----

```
sudo update-rc.d mysql defaults
```

-----*** Start mysql ***-----

```
sudo service mysql start
```

-----*** Mysql secure install ***----

```
$ sudo /usr/bin/mysql_secure_installation
[...]
Enter current password for root (enter for none):
OK, successfully used password, moving on...
[...]
Set root password? [Y/n] y
New password:
Re-enter new password:
Remove anonymous users? [Y/n] Y
[...]
Disallow root login remotely? [Y/n] N
[...]
Remove test database and access to it [Y/n] Y
[...]
Reload privilege tables now? [Y/n] Y
All done!
```

-----*** Install mysql JDBC driver ***-----

sudo apt-get install libmysql-java

Establish Your Cloudera Manager Repository

$ wget https://archive.cloudera.com/cm7/7.1.4/ubuntu1804/apt/cloudera-manager-trial.list

$ sudo mv cloudera-manager-trial.list /etc/apt/sources.list.d/

$ sudo apt-get update

$ sudo apt-key adv --recv-key --keyserver keyserver.ubuntu.com 73985D43B0B19C9F

$ sudo apt-get update

Install Cloudera Manager server software

---** Install Oracle JDK **---

$ sudo apt-get install openjdk-8-jdk -y

---** Install Cloudera Manager packages **---

$ sudo apt-get install cloudera-manager-daemons cloudera-manager-agent cloudera-manager-server -y

Setting up the Cloudera Manager Server Database

sudo mysql -u root

mysql>CREATE DATABASE cmdb;
mysql> CREATE USER 'cm'@'localhost' IDENTIFIED BY 'password';
mysql> GRANT ALL PRIVILEGES ON *.* TO 'cm'@'localhost' IDENTIFIED BY 'password' WITH GRANT OPTION;
mysql> exit;

sudo /opt/cloudera/cm/schema/scm_prepare_database.sh -p mysql cmdb cm password

Start the Cloudera Manager Server

$ sudo service cloudera-scm-server start

$ sudo service cloudera-scm-agent start

$ sudo service cloudera-scm-agent status

Login to CM web UI on 7180

After assigning service roles to hosts, when you are on database setup page, do following

Create database for Cloudera services

sudo mysql -u root

>> For reports manager

mysql> create database rman DEFAULT CHARACTER SET utf8;
mysql> CREATE USER 'rman'@'ip-172-31-30-152.us-west-1.compute.internal' IDENTIFIED BY 'password';
(Note : Enter pri-dns of host where the particular service is running)
mysql> grant all on rman.* TO 'rman'@'ip-172-31-30-152.us-west-1.compute.internal' IDENTIFIED BY 'password';

>> For Hive

mysql> create database metastore DEFAULT CHARACTER SET utf8;
mysql> CREATE USER 'hive'@'ip-172-31-27-181.us-west-1.compute.internal' IDENTIFIED BY 'password';
mysql> grant all on metastore.* TO 'hive'@'ip-172-31-27-181.us-west-1.compute.internal' IDENTIFIED BY 'password';

>> For Oozie

mysql> create database oozie default character set utf8;
mysql> CREATE USER 'oozie'@'ip-172-31-27-181.us-west-1.compute.internal' IDENTIFIED BY 'password';
mysql> grant all privileges on oozie.* to 'oozie'@'ip-172-31-27-181.us-west-1.compute.internal' identified by 'password';

== Add mysql driver jar

sudo cp /usr/share/java/mysql-connector-java-5.1.38.jar /opt/cloudera/parcels/CDH/lib/oozie/lib/

Install Mysql jdbc jar on hosts where services are running

sudo apt-get install libmysql-java

sudo cp /usr/share/java/mysql-connector-java-5.1.45.jar /opt/cloudera/parcels/CDH/lib/oozie/lib/

On database host

sudo su

cd /etc/mysql/mysql.conf.d/

sudo nano mysqld.cnf

bind-address=(private-ip-of-cm) OR (0.0.0.0)

$ sudo service mysql restart

Now continue installation on CM web ui

-----* Set up Edge node ***-----**

Add gateway roles on CM node

test whether read and write is working

Run a test job

yarn jar /opt/cloudera/parcels/CDH/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar pi 10 100

Set up Jumper node to access webui of daemons